

Tokenizing

By tokenizing, you can conveniently split up text by word or by sentence. This will allow you to work with smaller pieces of text that are still relatively coherent and meaningful even outside of the context of the rest of the text. It's your first step in turning unstructured data into structured data, which is easier to analyze.

When you're analyzing text, you'll be tokenizing by word and tokenizing by sentence. Here's what both types of tokenization bring to the table:

Tokenizing by word: Words are like the atoms of natural language. They're the smallest unit of meaning that still makes sense on its own. Tokenizing your text by word allows you to identify words that come up particularly often. For example, if you were analyzing a group of job ads, then you might find that the word "Python" comes up often. That could suggest high demand for Python knowledge, but you'd need to look deeper to know more.

Tokenizing by sentence: When you tokenize by sentence, you can analyze how those words relate to one another and see more context. Are there a lot of negative words around the word "Python" because the hiring manager doesn't like Python? Are there more terms from the domain of herpetology than the domain of software development, suggesting that you may be dealing with an entirely different kind of python than you were expecting?

Here's how to import the relevant parts of NLTK so you can tokenize by word and by sentence:

```
>>> from nltk.tokenize import sent_tokenize, word_tokenize
```

Now that you've imported what you need, you can create a string to tokenize. Here's a quote from Dune that you can use:

```
>>> example_string = """
```

```
... Muad'Dib learned rapidly because his first training was in how to learn.
```

... And the first lesson of all was the basic trust that he could learn.
... It's shocking to find how many people do not believe they can learn,
... and how many more believe learning to be difficult. """

You can use `sent_tokenize()` to split up `example_string` into sentences:

```
>>> sent_tokenize(example_string)
["Muad'Dib learned rapidly because his first training was in how to learn.",
 'And the first lesson of all was the basic trust that he could learn.',
 "It's shocking to find how many people do not believe they can learn, and how
 many more believe learning to be difficult."]
```

Tokenizing `example_string` by sentence gives you a list of three strings that are sentences:

1. "Muad'Dib learned rapidly because his first training was in how to learn."
2. 'And the first lesson of all was the basic trust that he could learn.'
3. "It's shocking to find how many people do not believe they can learn, and how many more believe learning to be difficult."

Now try tokenizing `example_string` by word:

```
>>> word_tokenize(example_string)
["Muad'Dib",
 'learned',
 'rapidly',
 'because',
 'his',
 'first',
```

'training',
'was',
'in',
'how',
'to',
'learn',
'',
'And',
'the',
'first',
'lesson',
'of',
'all',
'was',
'the',
'basic',
'trust',
'that',
'he',
'could',
'learn',
'',
'It',
'"s",
'shocking',
'to',
'find',

'how',
'many',
'people',
'do',
'not',
'believe',
'they',
'can',
'learn',
,,
'and',
'how',
'many',
'more',
'believe',
'learning',
'to',
'be',
'difficult',
'.]

You got a list of strings that NLTK considers to be words, such as:

- "Muad'Dib"
- 'training'
- 'how'

But the following strings were also considered to be words:

- "s"
- ','
- '.'

See how "It's" was split at the apostrophe to give you 'It' and "s", but "Muad'Dib" was left whole? This happened because NLTK knows that 'It' and "s" (a contraction of "is") are two distinct words, so it counted them separately. But "Muad'Dib" isn't an accepted contraction like "It's", so it wasn't read as two separate words and was left intact.

Filtering Stop Words

Stop words are words that you want to ignore, so you filter them out of your text when you're processing it. Very common words like 'in', 'is', and 'an' are often used as stop words since they don't add a lot of meaning to a text in and of themselves.

Here's how to import the relevant parts of NLTK in order to filter out stop words:

```
>>> nltk.download("stopwords")
>>> from nltk.corpus import stopwords
>>> from nltk.tokenize import word_tokenize
```

Here's a quote from Worf that you can filter:

```
>>> worf_quote = "Sir, I protest. I am not a merry man!"
```

Now tokenize worf_quote by word and store the resulting list in words_in_quote:

```
>>> words_in_quote = word_tokenize(worf_quote)
>>> words_in_quote
['Sir', ',', 'protest', ',', 'merry', 'man', '!']
```

You have a list of the words in `worf_quote`, so the next step is to create a set of stop words to filter `words_in_quote`. For this example, you'll need to focus on stop words in "english":

```
>>> stop_words = set(stopwords.words("english"))
```

Next, create an empty list to hold the words that make it past the filter:

```
>>> filtered_list = []
```

You created an empty list, `filtered_list`, to hold all the words in `words_in_quote` that aren't stop words. Now you can use `stop_words` to filter `words_in_quote`:

```
>>> for word in words_in_quote:
...     if word.casefold() not in stop_words:
...         filtered_list.append(word)
```

You iterated over `words_in_quote` with a for loop and added all the words that weren't stop words to `filtered_list`. You used `.casefold()` on `word` so you could ignore whether the letters in `word` were uppercase or lowercase. This is worth doing because `stopwords.words('english')` includes only lowercase versions of stop words.

Alternatively, you could use a list comprehension to make a list of all the words in your text that aren't stop words:

```
>>> filtered_list = [
...     word for word in words_in_quote if word.casefold() not in stop_words
... ]
```

When you use a list comprehension, you don't create an empty list and then add items to the end of it. Instead, you define the list and its contents at the same time. Using a list comprehension is often seen as more Pythonic.

Take a look at the words that ended up in `filtered_list`:

```
>>> filtered_list  
['Sir', ',', 'protest', '.', 'merry', 'man', '!']
```

Stemming

Stemming is a text processing task in which you reduce words to their root, which is the core part of a word. For example, the words “helping” and “helper” share the root “help.” Stemming allows you to zero in on the basic meaning of a word rather than all the details of how it’s being used. NLTK has more than one stemmer, but you’ll be using the Porter stemmer.

Here’s how to import the relevant parts of NLTK in order to start stemming:

```
>>> from nltk.stem import PorterStemmer  
>>> from nltk.tokenize import word_tokenize
```

Now that you’re done importing, you can create a stemmer with `PorterStemmer()`:

```
>>> stemmer = PorterStemmer()
```

The next step is for you to create a string to stem. Here’s one you can use:

```
>>> string_for_stemming = """  
... The crew of the USS Discovery discovered many discoveries.  
... Discovering is what explorers do."""
```

Before you can stem the words in that string, you need to separate all the words in it:

```
>>> words = word_tokenize(string_for_stemming)
```

Now that you have a list of all the tokenized words from the string, take a look at what's in words:

```
>>> words
```

```
['The',  
'crew',  
'of',  
'the',  
'USS',  
'Discovery',  
'discovered',  
'many',  
'discoveries',  
':',  
'Discovering',  
'is',  
'what',  
'explorers',  
'do',  
':']
```

Create a list of the stemmed versions of the words in words by using `stemmer.stem()` in a list comprehension:

```
>>> stemmed_words = [stemmer.stem(word) for word in words]
```

Take a look at what's in `stemmed_words`:

```
>>> stemmed_words
```

```
['the',
```


'crew',
'of',
'the',
'uss',
'discoveri',
'discov',
'mani',
'discoveri',
'.',
'discov',
'is',
'what',
'explor',
'do',
'.]

Here's what happened to all the words that started with 'discov' or 'Discov':

Original word	Stemmed version
'Discovery'	'discoveri'
'discovered'	'discov'
'discoveries'	'discoveri'
'Discovering'	'discov'

Those results look a little inconsistent. Why would 'Discovery' give you 'discoveri' when 'Discovering' gives you 'discov'?

Understemming and overstemming are two ways stemming can go wrong:

Understemming happens when two related words should be reduced to the same stem but aren't. This is a false negative.

Overstemming happens when two unrelated words are reduced to the same stem even though they shouldn't be. This is a false positive.

The Porter stemming algorithm dates from 1979, so it's a little on the older side. The Snowball stemmer, which is also called Porter2, is an improvement on the original and is also available through NLTK, so you can use that one in your own projects. It's also worth noting that the purpose of the Porter stemmer is not to produce complete words but to find variant forms of a word.

Fortunately, you have some other ways to reduce words to their core meaning, such as lemmatizing, which you'll see later in this tutorial. But first, we need to cover parts of speech.

Tagging Parts of Speech

Part of speech is a grammatical term that deals with the roles words play when you use them together in sentences. Tagging parts of speech, or POS tagging, is the task of labeling the words in your text according to their part of speech.

In English, there are eight parts of speech:

Part of speech	Role	Examples
Noun	Is a person, place, or thing	mountain, bagel, Poland
Pronoun	Replaces a noun	you, she, we
Adjective	Gives information about what a noun is like	efficient, windy, colorful
Verb	Is an action or a state of being	learn, is, go
Adverb	Gives information about a verb, an adjective, or another adverb	efficiently, always, very

Part of speech	Role	Examples
Preposition	Gives information about how a noun or pronoun is connected to another word	from, about, at
Conjunction	Connects two other words or phrases	so, because, and
Interjection	Is an exclamation	yay, ow, wow

Some sources also include the category articles (like “a” or “the”) in the list of parts of speech, but other sources consider them to be adjectives. NLTK uses the word determiner to refer to articles.

Here’s how to import the relevant parts of NLTK in order to tag parts of speech:

```
>>> from nltk.tokenize import word_tokenize
```

Now create some text to tag. You can use this Carl Sagan quote:

```
>>> sagan_quote = """
... If you wish to make an apple pie from scratch,
... you must first invent the universe."""
```

Use `word_tokenize` to separate the words in that string and store them in a list:

```
>>> words_in_sagan_quote = word_tokenize(sagan_quote)
```

Now call `nltk.pos_tag()` on your new list of words:

```
>>> import nltk
>>> nltk.pos_tag(words_in_sagan_quote)
[('If', 'IN'),
 ('you', 'PRP'),
```

('wish', 'VBP'),
('to', 'TO'),
('make', 'VB'),
('an', 'DT'),
('apple', 'NN'),
('pie', 'NN'),
('from', 'IN'),
('scratch', 'NN'),
(';', ';'),
('you', 'PRP'),
('must', 'MD'),
('first', 'VB'),
('invent', 'VB'),
('the', 'DT'),
('universe', 'NN'),
('.', '.')]]

Here's a summary that you can use to get started with NLTK's POS tags:

Tags that start with	Deal with
JJ	Adjectives
NN	Nouns
RB	Adverbs
PRP	Pronouns
VB	Verbs

Now that you know what the POS tags mean, you can see that your tagging was fairly successful:

- 'pie' was tagged NN because it's a singular noun.
- 'you' was tagged PRP because it's a personal pronoun.
- 'invent' was tagged VB because it's the base form of a verb.

But how would NLTK handle tagging the parts of speech in a text that is basically gibberish? Jabberwocky is a nonsense poem that doesn't technically mean much but is still written in a way that can convey some kind of meaning to English speakers.

Make a string to hold an excerpt from this poem:

```
>>> jabberwocky_excerpt = ""
```

```
... 'Twas brillig, and the slithy toves did gyre and gimble in the wabe:
```

```
... all mimsy were the borogoves, and the mome raths outgrabe.""
```

Use `word_tokenize` to separate the words in the excerpt and store them in a list:

```
>>> words_in_excerpt = word_tokenize(jabberwocky_excerpt)
```

Call `nltk.pos_tag()` on your new list of words:

```
>>> nltk.pos_tag(words_in_excerpt)
```

```
[('"'T', 'NN'),
```

```
 ('was', 'VBD'),
```

```
 ('brillig', 'VBN'),
```

```
 (',', ','),
```

```
 ('and', 'CC'),
```

```
 ('the', 'DT'),
```

('slithy', 'JJ'),
('toves', 'NNS'),
('did', 'VBD'),
('gyre', 'NN'),
('and', 'CC'),
('gimble', 'JJ'),
('in', 'IN'),
('the', 'DT'),
('wabe', 'NN'),
(':', ':'),
('all', 'DT'),
('mimsy', 'NNS'),
('were', 'VBD'),
('the', 'DT'),
('borogoves', 'NNS'),
(',', ','),
('and', 'CC'),
('the', 'DT'),
('mome', 'JJ'),
('raths', 'NNS'),
('outgrabe', 'RB'),
(':', ':')]

Accepted English words like 'and' and 'the' were correctly tagged as a conjunction and a determiner, respectively. The gibberish word 'slithy' was tagged as an adjective, which is what a human English speaker would probably assume from the context of the poem as well. Way to go, NLTK!

Lemmatizing

Now that you're up to speed on parts of speech, you can circle back to lemmatizing. Like stemming, lemmatizing reduces words to their core meaning, but it will give you a complete English word that makes sense on its own instead of just a fragment of a word like 'discoveri'.

Note: A lemma is a word that represents a whole group of words, and that group of words is called a lexeme.

For example, if you were to look up the word “blending” in a dictionary, then you’d need to look at the entry for “blend,” but you would find “blending” listed in that entry.

In this example, “blend” is the lemma, and “blending” is part of the lexeme. So when you lemmatize a word, you are reducing it to its lemma.

Here’s how to import the relevant parts of NLTK in order to start lemmatizing:

```
>>> from nltk.stem import WordNetLemmatizer
```

Create a lemmatizer to use:

```
>>> lemmatizer = WordNetLemmatizer()
```

Let’s start with lemmatizing a plural noun:

```
>>> lemmatizer.lemmatize("scarves")
```

```
'scarf'
```

"scarves" gave you 'scarf', so that’s already a bit more sophisticated than what you would have gotten with the Porter stemmer, which is 'scarv'. Next, create a string with more than one word to lemmatize:

```
>>> string_for_lemmatizing = "The friends of DeSoto love scarves."
```

Now tokenize that string by word:

```
>>> words = word_tokenize(string_for_lemmatizing)
```

Here's your list of words:

```
>>> words
```

```
['The',  
'friends',  
'of',  
'DeSoto',  
'love',  
'scarves',  
'']
```

Create a list containing all the words in words after they've been lemmatized:

```
>>> lemmatized_words = [lemmatizer.lemmatize(word) for word in words]
```

Here's the list you got:

```
>>> lemmatized_words
```

```
['The',  
'friend',  
'of',  
'DeSoto',  
'love',  
'scarf',  
'']
```


That looks right. The plurals 'friends' and 'scarves' became the singulars 'friend' and 'scarf'.

But what would happen if you lemmatized a word that looked very different from its lemma? Try lemmatizing "worst":

```
>>> lemmatizer.lemmatize("worst")
```

```
'worst'
```

You got the result 'worst' because `lemmatizer.lemmatize()` assumed that "worst" was a noun. You can make it clear that you want "worst" to be an adjective:

```
>>> lemmatizer.lemmatize("worst", pos="a")
```

```
'bad'
```

The default parameter for `pos` is 'n' for noun, but you made sure that "worst" was treated as an adjective by adding the parameter `pos="a"`. As a result, you got 'bad', which looks very different from your original word and is nothing like what you'd get if you were stemming. This is because "worst" is the superlative form of the adjective 'bad', and lemmatizing reduces superlatives as well as comparatives to their lemmas.

Now that you know how to use NLTK to tag parts of speech, you can try tagging your words before lemmatizing them to avoid mixing up homographs, or words that are spelled the same but have different meanings and can be different parts of speech.

Chunking

While tokenizing allows you to identify words and sentences, chunking allows you to identify phrases.

Note: A phrase is a word or group of words that works as a single unit to perform a grammatical function. Noun phrases are built around a noun.

Here are some examples:

“A planet”

“A tilting planet”

“A swiftly tilting planet”

Chunking makes use of POS tags to group words and apply chunk tags to those groups. Chunks don’t overlap, so one instance of a word can be in only one chunk at a time.

Here’s how to import the relevant parts of NLTK in order to chunk:

```
>>> from nltk.tokenize import word_tokenize
```

Before you can chunk, you need to make sure that the parts of speech in your text are tagged, so create a string for POS tagging. You can use this quote from The Lord of the Rings:

```
>>> lotr_quote = "It's a dangerous business, Frodo, going out your door."
```

Now tokenize that string by word:

```
>>> words_in_lotr_quote = word_tokenize(lotr_quote)
```

```
>>> words_in_lotr_quote
```

```
['It',  
 "'s",  
 'a',  
 'dangerous',  
 'business',  
 '']
```

```
'Frodo',  
,,  
'going',  
'out',  
'your',  
'door',  
'.]
```

Now you've got a list of all of the words in `lotr_quote`.

The next step is to tag those words by part of speech:

```
>>> nltk.download("averaged_perceptron_tagger")  
>>> lotr_pos_tags = nltk.pos_tag(words_in_lotr_quote)  
>>> lotr_pos_tags  
[('It', 'PRP'),  
 ('s', 'VBZ'),  
 ('a', 'DT'),  
 ('dangerous', 'JJ'),  
 ('business', 'NN'),  
 (',', ','),  
 ('Frodo', 'NNP'),  
 (',', ','),  
 ('going', 'VBG'),  
 ('out', 'RP'),  
 ('your', 'PRP$'),  
 ('door', 'NN'),  
 (',', ',')]
```

You've got a list of tuples of all the words in the quote, along with their POS tag. In order to chunk, you first need to define a chunk grammar.

Note: A chunk grammar is a combination of rules on how sentences should be chunked. It often uses regular expressions, or regexes.

For this tutorial, you don't need to know how regular expressions work, but they will definitely come in handy for you in the future if you want to process text.

Create a chunk grammar with one regular expression rule:

```
>>> grammar = "NP: {<DT>?<JJ>*<NN>}"
```

NP stands for noun phrase. You can learn more about noun phrase chunking in Chapter 7 of *Natural Language Processing with Python—Analyzing Text with the Natural Language Toolkit*.

According to the rule you created, your chunks:

1. Start with an optional (?) determiner ('DT')
2. Can have any number (*) of adjectives (JJ)
3. End with a noun (<NN>)

Create a chunk parser with this grammar:

```
>>> chunk_parser = nltk.RegexpParser(grammar)
```

Now try it out with your quote:

```
>>> tree = chunk_parser.parse(lotr_pos_tags)
```

Here's how you can see a visual representation of this tree:

```
>>> tree.draw()
```

This is what the visual representation looks like:

You got two noun phrases:

1. 'a dangerous business' has a determiner, an adjective, and a noun.
2. 'door' has just a noun.

Chinking

Chinking is used together with chunking, but while chunking is used to include a pattern, chinking is used to exclude a pattern.

Let's reuse the quote you used in the section on chunking. You already have a list of tuples containing each of the words in the quote along with its part of speech tag:

```
>>> lotr_pos_tags
```

```
[('It', 'PRP'),  
 ('s', 'VBZ'),  
 ('a', 'DT'),  
 ('dangerous', 'JJ'),  
 ('business', 'NN'),  
 (',', ','),  
 ('Frodo', 'NNP'),  
 (',', ','),  
 ('going', 'VBG'),  
 ('out', 'RP'),  
 ('your', 'PRP$'),  
 ('door', 'NN'),
```

```
('.', '.')]
```

The next step is to create a grammar to determine what you want to include and exclude in your chunks. This time, you're going to use more than one line because you're going to have more than one rule. Because you're using more than one line for the grammar, you'll be using triple quotes ("""):

```
>>> grammar = """
```

```
... Chunk: {<.*>+}
```

```
...     }<JJ>{"""
```

The first rule of your grammar is {<.*>+}. This rule has curly braces that face inward ({}) because it's used to determine what patterns you want to include in your chunks. In this case, you want to include everything: <.*>+.

The second rule of your grammar is }<JJ>{. This rule has curly braces that face outward ({}) because it's used to determine what patterns you want to exclude in your chunks. In this case, you want to exclude adjectives: <JJ>.

Create a chunk parser with this grammar:

```
>>> chunk_parser = nltk.RegexpParser(grammar)
```

Now chunk your sentence with the chunk you specified:

```
>>> tree = chunk_parser.parse(lotr_pos_tags)
```

You get this tree as a result:

```
>>> tree
```

```
Tree('S', [Tree('Chunk', [(('It', 'PRP'), ('s', 'VBZ'), ('a', 'DT'))], ('dangerous', 'JJ'),  
Tree('Chunk', [(('business', 'NN'), (',', ','), ('Frodo', 'NNP'), (',', ','), ('going', 'VBG'),  
(('out', 'RP'), ('your', 'PRP$'), ('door', 'NN'), (',', ','))]))])
```

In this case, ('dangerous', 'JJ') was excluded from the chunks because it's an adjective (JJ). But that will be easier to see if you get a graphic representation again:

Here, you've excluded the adjective 'dangerous' from your chunks and are left with two chunks containing everything else. The first chunk has all the text that appeared before the adjective that was excluded. The second chunk contains everything after the adjective that was excluded.

Using Named Entity Recognition (NER)

Named entities are noun phrases that refer to specific locations, people, organizations, and so on. With named entity recognition, you can find the named entities in your texts and also determine what kind of named entity they are.

Here's the list of named entity types from the NLTK book:

NE type	Examples
ORGANIZATION	Georgia-Pacific Corp., WHO
PERSON	Eddy Bonte, President Obama
LOCATION	Murray River, Mount Everest
DATE	June, 2008-06-29
TIME	two fifty a m, 1:30 p.m.
MONEY	175 million Canadian dollars, GBP 10.40
PERCENT	twenty pct, 18.75 %
FACILITY	Washington Monument, Stonehenge
GPE	South East Asia, Midlothian

You can use `nltk.ne_chunk()` to recognize named entities. Let's use `lotr_pos_tags` again to test it out:

```
>>> nltk.download("maxent_ne_chunker")
>>> nltk.download("words")
>>> tree = nltk.ne_chunk(lotr_pos_tags)
```

That's how you can identify named entities! But you can take this one step further and extract named entities directly from your text. Create a string from which to extract named entities. You can use this quote from *The War of the Worlds*:

```
>>> quote = """
... Men like Schiaparelli watched the red planet—it is odd, by-the-bye, that
... for countless centuries Mars has been the star of war—but failed to
... interpret the fluctuating appearances of the markings they mapped so well.
... All that time the Martians must have been getting ready.
...
... During the opposition of 1894 a great light was seen on the illuminated
... part of the disk, first at the Lick Observatory, then by Perrotin of Nice,
... and then by other observers. English readers heard of it first in the
... issue of Nature dated August 2."""
```

Now create a function to extract named entities:

```
>>> def extract_ne(quote):
...     words = word_tokenize(quote, language=language)
...     tags = nltk.pos_tag(words)
...     tree = nltk.ne_chunk(tags, binary=True)
...     return set(
...         " ".join(i[0] for i in t)
...         for t in tree
```



```
...     if hasattr(t, "label") and t.label() == "NE"
... )
```

With this function, you gather all named entities, with no repeats. In order to do that, you tokenize by word, apply part of speech tags to those words, and then extract named entities based on those tags. Because you included `binary=True`, the named entities you'll get won't be labeled more specifically. You'll just know that they're named entities.

Take a look at the information you extracted:

```
>>> extract_ne(quote)
{'Lick Observatory', 'Mars', 'Nature', 'Perrotin', 'Schiaparelli'}
```

You missed the city of Nice, possibly because NLTK interpreted it as a regular English adjective, but you still got the following:

- An institution: 'Lick Observatory'
- A planet: 'Mars'
- A publication: 'Nature'
- People: 'Perrotin', 'Schiaparelli'

Getting Text to Analyze

Now that you've done some text processing tasks with small example texts, you're ready to analyze a bunch of texts at once. A group of texts is called a corpus. NLTK provides several corpora covering everything from novels hosted by Project Gutenberg to inaugural speeches by presidents of the United States.

In order to analyze texts in NLTK, you first need to import them. This requires `nltk.download("book")`, which is a pretty big download:

```
>>> nltk.download("book")
>>> from nltk.book import *
*** Introductory Examples for the NLTK Book ***
```

Loading text1, ..., text9 and sent1, ..., sent9

Type the name of the text or sentence to view it.

Type: 'texts()' or 'sents()' to list the materials.

text1: Moby Dick by Herman Melville 1851

text2: Sense and Sensibility by Jane Austen 1811

text3: The Book of Genesis

text4: Inaugural Address Corpus

text5: Chat Corpus

text6: Monty Python and the Holy Grail

text7: Wall Street Journal

text8: Personals Corpus

text9: The Man Who Was Thursday by G . K . Chesterton 1908

You now have access to a few linear texts (such as Sense and Sensibility and Monty Python and the Holy Grail) as well as a few groups of texts (such as a chat corpus and a personals corpus). Human nature is fascinating, so let's see what we can find out by taking a closer look at the personals corpus!

This corpus is a collection of personals ads, which were an early version of online dating. If you wanted to meet someone, then you could place an ad in a newspaper and wait for other readers to respond to you.

If you'd like to learn how to get other texts to analyze, then you can check out Chapter 3 of Natural Language Processing with Python – Analyzing Text with the Natural Language Toolkit.

Using a Concordance

When you use a concordance, you can see each time a word is used, along with its immediate context. This can give you a peek into how a word is being used at the sentence level and what words are used with it.

Let's see what these good people looking for love have to say! The personals corpus is called text8, so we're going to call `.concordance()` on it with the parameter "man":

```
>>> text8.concordance("man")
```

Displaying 14 of 14 matches:

to hearing from you all . ABLE young man seeks , sexy older women . Phone for
ble relationship . GENUINE ATTRACTIVE MAN 40 y . o . , no ties , secure , 5 ft .
ship , and quality times . VIETNAMESE MAN Single , never married , financially
ip . WELL DRESSED emotionally healthy man 37 like to meet full figured woman fo
nth subs LIKE TO BE MISTRESS of YOUR MAN like to be treated well . Bold DTE
no
eeks lady in similar position MARRIED MAN 50 , attrac . fit , seeks lady 40 - 5
eks nice girl 25 - 30 serious rship . Man 46 attractive fit , assertive , and k
40 - 50 sought by Aussie mid 40s b / man f / ship r / ship LOVE to meet widowe
discreet times . Sth E Subs . MARRIED MAN 42yo 6ft , fit , seeks Lady for discr
woman , seeks professional , employed man , with interests in theatre , dining
tall and of large build seeks a good man . I am a nonsmoker , social drinker ,
lead to relationship . SEEKING HONEST MAN I am 41 y . o . , 5 ft . 4 , med . bui
quiet times . Seeks 35 - 45 , honest man with good SOH & similar interests , f
genuine , caring , honest and normal man for fship , poss rship . S / S , S /

Interestingly, the last three of those fourteen matches have to do with seeking an honest man, specifically:

1. SEEKING HONEST MAN
2. Seeks 35 - 45 , honest man with good SOH & similar interests
3. genuine , caring , honest and normal man for fship , poss rship

Let's see if there's a similar pattern with the word "woman":

```
>>> text8.concordance("woman")
```

Displaying 11 of 11 matches:

at home . Seeking an honest , caring woman , slim or med . build , who enjoys t
thy man 37 like to meet full figured woman for relationship . 48 slim , shy , S
rry . MALE 58 years old . Is there a Woman who would like to spend 1 weekend a
other interests . Seeking Christian Woman for fship , view to rship . SWM 45 D
ALE 60 - burly beared seeks intimate woman for outings n / s s / d F / ston / P
ington . SCORPIO 47 seeks passionate woman for discreet intimate encounters
SEX

le dad . 42 , East sub . 5 " 9 seeks woman 30 + for f / ship relationship TALL
personal trainer looking for married woman age open for fun MARRIED Dark guy
37

rinker , seeking slim - medium build woman who is happy in life , age open . AC
. O . TERTIARY Educated professional woman , seeks professional , employed
man

real romantic , age 50 - 65 y . o . WOMAN OF SUBSTANCE 56 , 59 kg . , 50 , fit

The issue of honesty came up in the first match only:

Seeking an honest , caring woman , slim or med . build

Dipping into a corpus with a concordance won't give you the full picture, but it can still be interesting to take a peek and see if anything stands out.

Making a Dispersion Plot

You can use a dispersion plot to see how much a particular word appears and where it appears. So far, we've looked for "man" and "woman", but it would be interesting to see how much those words are used compared to their synonyms:

```
>>> text8.dispersion_plot(  
...    ["woman", "lady", "girl", "gal", "man", "gentleman", "boy", "guy"]  
... )
```

Each vertical blue line represents one instance of a word. Each horizontal row of blue lines represents the corpus as a whole. This plot shows that:

"lady" was used a lot more than "woman" or "girl". There were no instances of "gal".

"man" and "guy" were used a similar number of times and were more common than "gentleman" or "boy".

You use a dispersion plot when you want to see where words show up in a text or corpus. If you're analyzing a single text, this can help you see which words show up near each other. If you're analyzing a corpus of texts that is organized chronologically, it can help you see which words were being used more or less over a period of time.

Staying on the theme of romance, see what you can find out by making a dispersion plot for *Sense and Sensibility*, which is text2. Jane Austen novels talk a lot about people's homes, so make a dispersion plot with the names of a few homes:

```
>>> text2.dispersion_plot(["Allenham", "Whitwell", "Cleveland", "Combe"])
```

Apparently Allenham is mentioned a lot in the first third of the novel and then doesn't come up much again. Cleveland, on the other hand, barely comes up in the first two thirds but shows up a fair bit in the last third. This distribution reflects changes in the relationship between Marianne and Willoughby:

- Allenham is the home of Willoughby's benefactress and comes up a lot when Marianne is first interested in him.
- Cleveland is a home that Marianne stays at after she goes to see Willoughby in London and things go wrong.

Dispersion plots are just one type of visualization you can make for textual data. The next one you'll take a look at is frequency distributions.

Making a Frequency Distribution

With a frequency distribution, you can check which words show up most frequently in your text. You'll need to get started with an import:

```
>>> from nltk import FreqDist
```

FreqDist is a subclass of `collections.Counter`. Here's how to create a frequency distribution of the entire corpus of personals ads:

```
>>> frequency_distribution = FreqDist(text8)
```

```
>>> print(frequency_distribution)
```

```
<FreqDist with 1108 samples and 4867 outcomes>
```

Since 1108 samples and 4867 outcomes is a lot of information, start by narrowing that down. Here's how to see the 20 most common words in the corpus:

```
>>> frequency_distribution.most_common(20)
```

```
[(',', 539),
```

```
('.', 353),
```

```
('/', 110),
```

```
('for', 99),
```

```
('and', 74),
```

```
('to', 74),
```

```
('lady', 68),
```

```
('-', 66),
```

```
('seeks', 60),
```

```
('a', 52),
```

```
('with', 44),  
('S', 36),  
('ship', 33),  
('&', 30),  
('relationship', 29),  
('fun', 28),  
('in', 27),  
('slim', 27),  
('build', 27),  
('o', 26)]
```

You have a lot of stop words in your frequency distribution, but you can remove them just as you did earlier. Create a list of all of the words in text8 that aren't stop words:

```
>>> meaningful_words = [  
...     word for word in text8 if word.casefold() not in stop_words  
... ]
```

Now that you have a list of all of the words in your corpus that aren't stop words, make a frequency distribution:

```
>>> frequency_distribution = FreqDist(meaningful_words)
```

Take a look at the 20 most common words:

```
>>> frequency_distribution.most_common(20)  
[(',', 539),  
('!', 353),  
('/', 110),  
('lady', 68),
```

('-', 66),
('seeks', 60),
('ship', 33),
('&', 30),
('relationship', 29),
('fun', 28),
('slim', 27),
('build', 27),
('smoker', 23),
('50', 23),
('non', 22),
('movies', 22),
('good', 21),
('honest', 20),
('dining', 19),
('rship', 18)]

Some of the most common words are:

- 'lady'
- 'seeks'
- 'ship'
- 'relationship'
- 'fun'
- 'slim'
- 'build'
- 'smoker'
- '50'
- 'non'
- 'movies'

- 'good'
- 'honest'

Finding Collocations

A collocation is a sequence of words that shows up often. If you're interested in common collocations in English, then you can check out The BBI Dictionary of English Word Combinations. It's a handy reference you can use to help you make sure your writing is idiomatic. Here are some examples of collocations that use the word "tree":

- Syntax tree
- Family tree
- Decision tree

To see pairs of words that come up often in your corpus, you need to call `.collocations()` on it:

```
>>> text8.collocations()
```

```
would like; medium build; social drinker; quiet nights; non smoker;  
long term; age open; Would like; easy going; financially secure; fun  
times; similar interests; Age open; weekends away; poss rship; well  
presented; never married; single mum; permanent relationship; slim  
build
```

slim build did show up, as did medium build and several other word combinations. No long walks on the beach though!

But what would happen if you looked for collocations after lemmatizing the words in your corpus? Would you find some word combinations that you missed the first time around because they came up in slightly varied versions?

If you followed the instructions earlier, then you'll already have a lemmatizer, but you can't call `collocations()` on just any data type, so you're going to need to do

some prep work. Start by creating a list of the lemmatized versions of all the words in text8:

```
>>> lemmatized_words = [lemmatizer.lemmatize(word) for word in text8]
```

But in order for you to be able to do the linguistic processing tasks you've seen so far, you need to make an NLTK text with this list:

```
>>> new_text = nltk.Text(lemmatized_words)
```

Here's how to see the collocations in your new_text:

```
>>> new_text.collocations()
```

medium build; social drinker; non smoker; long term; would like; age
open; easy going; financially secure; Would like; quiet night; Age
open; well presented; never married; single mum; permanent
relationship; slim build; year old; similar interest; fun time; Photo
pls

Compared to your previous list of collocations, this new one is missing a few:

- weekends away
- poss rship

The idea of quiet nights still shows up in the lemmatized version, quiet night. Your latest search for collocations also brought up a few news ones:

- year old suggests that users often mention ages.
- photo pls suggests that users often request one or more photos.